

Statistical Thinking in Python Part 1

3). Thinking probabilistically-- Discrete variables

a). Generating random numbers using the np.random module

Seed the random number generator

```
np.random.seed(42)
```

Initialize random numbers: random_numbers

```
random_numbers = np.empty(100000)
```

Generate random numbers by looping over range(100000)

```
for i in range(100000):
```


```
    random_numbers[i] = np.random.random()
```

Plot a histogram

```
_ = plt.hist(random_numbers)
```

Show the plot

```
plt.show()
```



b). The np.random module and Bernoulli trials

```
def perform_bernoulli_trials(n, p):  
    """Perform n Bernoulli trials with success probability p  
    and return number of successes."""  
    # Initialize number of successes: n_success  
    n_success = 0  
  
    # Perform trials  
    for i in range(n):  
        # Choose random number between zero and one: random_number  
        random_number = np.random.random()  
  
        # If less than p, it's a success so add one to n_success  
        if random_number < p:  
            n_success += 1  
  
    return n_success
```

c). How many defaults we might expect

```
# Seed random number generator
```

```
np.random.seed(42)
```

```
# Initialize the number of defaults: n_defaults
```

```
n_defaults= np.empty(1000)
```

```
# Compute the number of defaults
```

```
for i in range(1000):
```

```
    n_defaults[i] = perform_bernoulli_trials(100,0.05)
```

```
# Plot the histogram with default number of bins; label your axes
```


```
_ = plt.hist(n_defaults, normed=True)
```

```
_ = plt.xlabel('number of defaults out of 100 loans')
```

```
_ = plt.ylabel('probability')
```

```
# Show the plot
```

```
plt.show()
```



d). Will the bank fail

```
# Compute ECDF: x, y
```

```
x,y= ecdf(n_defaults)
```

```
# Plot the ECDF with labeled axes
```

```
plt.plot(x, y, marker='.', linestyle='none')
```

```
plt.xlabel('x')
```

```
plt.ylabel('y')
```

```
# Show the plot
```


```
plt.show()
```

```
# Compute the number of 100-loan simulations with 10 or more defaults: n_lose_money
```

```
n_lose_money=np.sum(n_defaults >= 10)
```

```
# Compute and print probability of losing money
```

```
print('Probability of losing money =', n_lose_money / len(n_defaults))
```



e). Sampling out of the Binomial Distribution

```
# Take 10,000 samples out of the binomial distribution: n_defaults
```

```
n_defaults=np.random.binomial(n=100, p=0.05, size=10000)
```

```
# Compute CDF: x, y
```

```
x,y=ecdf(n_defaults)
```

```
# Plot the CDF with axis labels
```

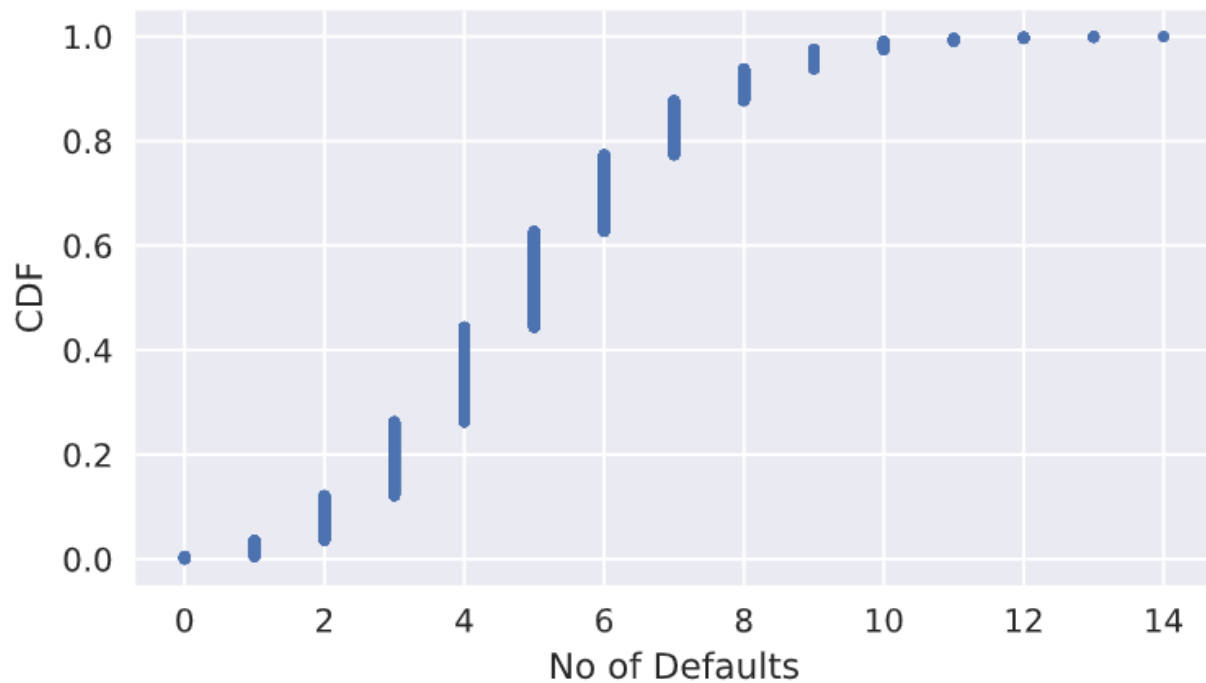
```
_=plt.plot(x, y, marker='.', linestyle='none')
```

```
_=plt.xlabel('No of Defaults')
```

```
_=plt.ylabel('CDF')
```

```
# Show the plot
```

```
plt.show()
```



f). Plotting the binomial PMF

```
# Compute bin edges: bins
```

```
bins = np.arange(min(n_defaults), max(n_defaults) + 1.5) - 0.5
```

```
# Generate histogram
```

```
plt.hist(n_defaults, normed=True, bins=bins)
```


```
# Label axes
```

```
plt.xlabel('x')
```

```
plt.ylabel('y')
```

```
# Show the plot
```

```
plt.show()
```



g). Relationship between Binomial and Poisson Distribution:

```
# Draw 10,000 samples out of Poisson distribution: samples_poisson
samples_poisson= np.random.poisson(10, size= 10000)

# Print the mean and standard deviation
print('Poisson:  ', np.mean(samples_poisson), np.std(samples_poisson))

# Specify values of n and p to consider for Binomial: n, p
n=[20,100,1000]
p=[0.5,0.1,0.01]

# Draw 10,000 samples for each n,p pair: samples_binomial
for i in range(3):
    samples_binomial = np.random.binomial(n[i], p[i],10000)

# Print results
print('n =', n[i], 'Binom:', np.mean(samples_binomial), np.std(samples_binomial))

<script.py> output:
Poisson:   10.0186 3.144813832327758
n = 20 Binom: 9.9637 2.2163443572694206
n = 100 Binom: 9.9947 3.0135812433050484
n = 1000 Binom: 9.9985 3.139378561116833
```

h). Was 2015 anomalous**# Draw 10,000 samples out of Poisson distribution: n_nohitters****n_nohitters= np.random.poisson(251/115, size=10000)****# Compute number of samples that are seven or greater: n_large****n_large = np.sum(n_nohitters >= 7)****# Compute probability of getting seven or more: p_large****p_large= n_large/10000****# Print the result****print('Probability of seven or more no-hitters:', p_large)****<script.py> output:****Probability of seven or more no-hitters: 0.0067**